


Operating Systems Programming Assignment #7

File Fragmentation Reproduction

Prof. Li-Pin Chang, NCTU

File fragmentation

- Ideally, a file is allocated in contiguous disk space
 - File system ages after it undergoes a number of allocation and de-allocation requests
 - Free space become fragmented
 - It is difficult for an aged file system to allocate contiguous disk space for a new file or an existing to grow
- File fragmentation increase the disk positioning overhead



Extents

- An extent = multiple contiguous disk blocks
 - Ext4 allocates in terms of extents to avoid file fragmentation
 - With delayed allocation and pre-allocation
- You can use <filefrag -v filename> command to get the information of the extents of a file

```
jm9935333@ubuntu:~/newdisk$ filefrag largefile.txt
largefile.txt: 275 extents found
jm9935333@ubuntu:~/newdisk$ filefrag -v largefile.txt
Filesystem type is: ef53
File size of largefile.txt is 44539904 (43496 blocks of 1024 bytes)
ext:      logical_offset:      physical_offset: length:      expected: flags:
 0:         0..      2047:         6145..      8192:      2048:
 1:      2048..      4095:         4030..      6077:      2048:      8193:
 2:      4096..      4349:         57603..     57856:      254:      6078:
 3:      4350..      4603:         73987..     74240:      254:      57857:
 4:      4604..      4729:         41219..     41344:      126:      74241:
 5:      4730..      4809:         102321..    102400:       80:      41345:
```

APIs & Tools

- `fopen()`, `fclose()`, `fread()`, `fwrite()`, `fsync()`, etc
- `filefrag [-Bbsvx] [files.....]`
 - Reports the details of a file's extents
 - <https://linux.die.net/man/8/filefrag>

Procedure

- Your program
 - gets time t1
 - runs on a 100 MB disk volume formatted in ext4
 - ages the 100 MB file system with your own method
 - produces a file “largefile.txt”
 - `system(“filefrag -v largefile.txt”);`
 - gets time t2
 - prints t2-t1 (in seconds)

Requirements

- “largefile.txt” must have ≥ 400 extents
- The virtual disk contains only your program and files created by it
 - The sizes of created new files must be powers of 2
- Your program must not run too slow, see the grading policy below

Test Flow

1. Attach a 100 MB virtual disk to your VM
2. Format the disk in ext4
3. Mount the disk
4. Copy your program to the new volume
5. Run your program
 - Age the file system with your own method...
 - Create “largefile.txt”
 - Call `system(“filefrag -v largefile.txt”)`
 - Print elapsed time in xxx.xx seconds

Grading

- 30% meet all the requirements
 - Program flow, file name “largfile.txt”, file sizes (powers of 2)...
- 30% extent count of “largefile.txt”
 - 30 pts if ≥ 400
 - 20 pts if in $[200, 400]$
 - 0 pts if < 200
- 40% program efficiency
 - Let X =the median execution time of all students' programs
 - Let Y =the execution of your program, you will get
 - 40 pts if $Y \leq X$
 - 35 pts if Y is in $(X, 2X]$
 - 20 pts if $Y > 2X$
- Extent # and execution times will be measured on the TA's platform

Example

```
Filesystem type is: ef53
File size of /home/charlie/文件/HW/100MBext4/largefile.txt is 44620800 (43575 blocks of 1024 bytes)
```

ext:	logical_offset:	physical_offset:	length:	expected:	flags:
0:	0..	2047:	6145..	8192:	2048:
1:	2048..	4095:	3584..	5631:	2048:
2:	4096..	6128:	100368..	102400:	2033:
3:	6129..	6143:	103409..	103423:	15:
4:	6144..	6656:	5632..	6144:	513:
5:	6657..	6718:	24835..	24896:	62:
6:	6719..	6799:	41219..	41299:	62:

...

```
579: 43520.. 43574: 24065.. 24119: 55: 24001: last,eof
/home/charlie/文件/HW/100MBext4/largefile.txt: 580 extents found
Elapsed time: 19.983069 sec
```

Testing OS Environment

- Ubuntu 16.04, Ubuntu 14.04 or CS linux work station
 - Your code should compile successfully in one of the above environments